


# Probabilistic Supervised Learning: Linear Regression

CS772A: Probabilistic Machine Learning

Piyush Rai

# Probabilistic Supervised Learning

- Goal: To learn the conditional distribution  $p(y|x)$  of output given input
- The form of the distribution  $p(y|x)$  depends on output type, e.g.,
  - Real: Model  $p(y|x)$  using a Gaussian (or some other suitable real-valued distribution)
  - Binary: Model  $p(y|x)$  using a Bernoulli
  - Categorical/multiclass: Model  $p(y|x)$  using a multinoulli/categorical distribution
  - Various other types (e.g., count, positive reals, etc) can also be modeled using appropriate distributions (e.g., Poisson for count, gamma for positive reals)
- The distribution  $p(y|x)$  can be defined directly or indirectly



“Direct” way without modeling the inputs  $x_n$

Parameters of this distribution are the outputs of function  $f$

“Indirect” way by modeling the outputs as well as the inputs

$$p(y|x) = p(y|f(x, w))$$

$$p(y|x) = \frac{p(y, x)}{p(x)}$$

“Indirect” way requires first learning the joint distribution of inputs and outputs



# Discriminative vs Generative Sup. Learning

Non-probabilistic supervised learning approaches (e.g., SVM) are usually considered discriminative since  $p(\mathbf{x})$  is never modeled

- Direct way of sup. learning is discriminative, indirect way is generative

## Discriminative Approach

$$p(y|\mathbf{x}) = p(y|f(\mathbf{x}, \mathbf{w}))$$

$f$  can be any function which uses inputs and weights  $\mathbf{w}$  to defines parameters of distr.  $p$

Some examples

$$p(y|\mathbf{x}) = \mathcal{N}(y|\mathbf{w}^\top \mathbf{x}, \beta^{-1})$$

$$p(y|\mathbf{x}) = \text{Bernoulli}(y|\sigma(\mathbf{w}^\top \mathbf{x}))$$

## Generative Approach

$$p(y|\mathbf{x}) = \frac{p(y, \mathbf{x})}{p(\mathbf{x})}$$

Requires estimating the **joint distribution** of inputs and outputs to get the conditional  $p(y|\mathbf{x})$  (unlike the discriminative approach which directly estimates the conditional  $p(y|\mathbf{x})$  and does not model the distribution of  $\mathbf{x}$ )

- Note: Generative approach can also be used for other settings too, such as unsupervised learning and semi-supervised learning (will see later)



# Probabilistic Linear Regression

A discriminative model for regression problems

- Assume training data  $\{\mathbf{x}_n, \mathbf{y}_n\}_{n=1}^N$ , with features  $\mathbf{x}_n \in \mathbb{R}^D$  and responses  $\mathbf{y}_n \in \mathbb{R}$ 
  - Unknown to be estimated
- Assume  $\mathbf{y}_n$  generated by a noisy linear model with wts  $\mathbf{w} = [w_1, \dots, w_D] \in \mathbb{R}^D$ 
  - Each weight assumed real-valued



$$\mathbf{y}_n = \mathbf{w}^\top \mathbf{x}_n + \epsilon_n$$

Gaussian noise drawn from  $\mathcal{N}(\epsilon_n | 0, \beta^{-1})$

- Notation alert:  $\beta$  is the precision (and  $\beta^{-1}$  the variance)
  - Unknown to be estimated

## Likelihood model

$$p(\mathbf{y}_n | \mathbf{x}_n, \mathbf{w}, \beta) = \mathcal{N}(\mathbf{y}_n | \mathbf{w}^\top \mathbf{x}_n, \beta^{-1})$$



# Probabilistic Linear Regression

- For all the training data, we can write the above model in matrix-vector notation

$\mathbf{y} = [y_1; y_2; \dots; y_N]$  is the  $N \times 1$  response vector

$\mathbf{X} = [\mathbf{x}_1^\top; \mathbf{x}_2^\top; \dots; \mathbf{x}_N^\top]$  is the  $N \times D$  input matrix


$\boldsymbol{\epsilon} = [\epsilon_1; \epsilon_2; \dots; \epsilon_N]$  is the  $N \times 1$  noise vector drawn from  $\mathcal{N}(\mathbf{0}, \beta^{-1} \mathbf{I}_N)$

Same as writing

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w}, \beta) = \mathcal{N}(\mathbf{y}|\mathbf{X}\mathbf{w}, \beta^{-1} \mathbf{I}_N)$$

$$\mathbf{y} = \mathbf{X}\mathbf{w} + \boldsymbol{\epsilon}$$

- This is a linear Gaussian model with  $\mathbf{w}$  being the unknown Gaussian r.v.
- A simple “plate diagram” for this model would look like this (hyperparameters not shown in the diagram)



# On compact notations..

- When writing the likelihood (assuming  $\mathbf{y}_n$ 's are i.i.d. given  $\mathbf{w}$  and  $\mathbf{x}_n$ )

$$\begin{aligned} p(\mathbf{y}|\mathbf{X}, \mathbf{w}, \beta) &= \prod_{n=1}^N \mathcal{N}(y_n | \mathbf{w}^\top \mathbf{x}_n, \beta^{-1}) \\ &= \mathcal{N}(\mathbf{y} | \mathbf{X}\mathbf{w}, \beta^{-1} \mathbf{I}_N) \end{aligned}$$

- Thus a product of  $N$  univariate Gaussians here (not always) is equivalent to an  $N$ -dim Gaussian over the vector  $\mathbf{y} = [y_1, y_2, \dots, y_N]$
- We will prefer to use this equivalence at other places too whenever we have multiple i.i.d. random variables, each having a univariate Gaussian distribution



# Prior on weights


- Assume a **zero-mean Gaussian prior** on  $\mathbf{w}$

$$p(\mathbf{w}|\lambda) = \prod_{d=1}^D p(w_d|\lambda) = \prod_{d=1}^D \mathcal{N}(w_d|0, \lambda^{-1})$$

In zero-mean case,  $\lambda$  sort of denotes each feature's importance. Think why?

Large  $\lambda$  means more aggressive push towards zero

The precision  $\lambda$  controls how aggressively the prior pushes  $w_d$  towards mean (0)



This prior assumes that *a priori* each weight has a small value (close to zero)

$\lambda$  controls the uncertainty around our prior belief about value of  $w_d$

$$= \mathcal{N}(\mathbf{w}|\mathbf{0}, \lambda^{-1} \mathbf{I}_D)$$

$$\propto \left(\frac{\lambda}{2\pi}\right)^{\frac{D}{2}} \exp\left[-\frac{\lambda}{2} \mathbf{w}^\top \mathbf{w}\right]$$

May also use a non-zero mean Gaussian prior, e.g.,  $\mathcal{N}(\mathbf{w}_d|\mu, \lambda^{-1})$  if we expect weights to be close to some value  $\mu$



Can also use a **full covariance matrix**  $\Lambda^{-1}$  for the prior to impose a priori correlations among different weights

Prior's hyperparameters ( $\lambda/\Lambda/\mu$ ) etc can be learned as well using point estimation (e.g., MLE-II) or fully Bayesian inference

- Zero-mean Gaussian prior corresponds to  $\ell_2$  regularizer

Reason: The negative log prior  $-\log p(\mathbf{w}) \propto \frac{\lambda}{2} \mathbf{w}^\top \mathbf{w}$





# The Posterior

MLE/MAP left  
as an exercise



- The posterior over  $\mathbf{w}$  (for now, assume hyperparams  $\beta$  and  $\lambda$  to be known)

$$p(\mathbf{w}|\mathbf{y}, \mathbf{X}, \beta, \lambda) = \frac{p(\mathbf{w}|\lambda)p(\mathbf{y}|\mathbf{w}, \mathbf{X}, \beta)}{p(\mathbf{y}|\mathbf{X}, \beta, \lambda)} \propto p(\mathbf{w}|\lambda)p(\mathbf{y}|\mathbf{w}, \mathbf{X}, \beta)$$

Must be a Gaussian  
due to conjugacy

Marginal likelihood for this regression model.  
Note that it is conditioned on  $\mathbf{X}$  too which is  
assumed given and not being modeled

$$p(\mathbf{w}|\mathbf{y}, \mathbf{X}, \beta, \lambda) \propto \mathcal{N}(\mathbf{w}|\mathbf{0}, \lambda^{-1}\mathbf{I}_D) \times \mathcal{N}(\mathbf{y}|\mathbf{X}\mathbf{w}, \beta^{-1}\mathbf{I}_N)$$

- Using the “completing the squares” trick (or linear Gaussian model results)

$$p(\mathbf{w}|\mathbf{y}, \mathbf{X}, \beta, \lambda) = \mathcal{N}(\mu_N, \Sigma_N)$$

Note that  $\lambda$  and  $\beta$  can be  
learned under the  
probabilistic set-up (though  
assumed fixed as of now)

where  $\Sigma_N = (\beta \sum_{n=1}^N \mathbf{x}_n \mathbf{x}_n^\top + \lambda \mathbf{I}_D)^{-1} = (\beta \mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I}_D)^{-1}$  (posterior's covariance matrix)

The form is also similar to the solution to ridge regression  
 $\text{argmin}_{\mathbf{w}} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|^2 + \lambda \mathbf{w}^\top \mathbf{w} = (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^\top \mathbf{y}$

MAP solution turns out to be exactly  
the same (reason: Gaussian's mean  
and mode are the same)


$$\mu_N = \Sigma_N \left[ \beta \sum_{n=1}^N y_n \mathbf{x}_n \right] = \Sigma_N [\beta \mathbf{X}^\top \mathbf{y}] = (\mathbf{X}^\top \mathbf{X} + \frac{\lambda}{\beta} \mathbf{I}_D)^{-1} \mathbf{X}^\top \mathbf{y} \quad (\text{posterior's mean})$$



# The Posterior: A Visualization

- Assume a lin. reg. problem with true  $\mathbf{w} = [w_0, w_1]$ ,  $w_0 = -0.3, w_1 = 0.5$
- Assume data generated by a linear regression model  $y = w_0 + w_1x + \text{"noise"}$ 
  - Note: It's actually 1-D regression ( $w_0$  is just a bias term), or 2-D reg. with feature  $[1, x]$
- Figures below show the “data space” and posterior of  $\mathbf{w}$  for different number of observations (note: with no observations, the posterior = prior)

Each red line represents the “data” generated for a randomly drawn  $\mathbf{w}$  from the current posterior



# Posterior Predictive Distribution

- To get the prediction  $y_*$  for a new input  $\mathbf{x}_*$ , we can compute its PPD

$$p(y_* | \mathbf{x}_*, \mathbf{X}, \mathbf{y}, \beta, \lambda) = \int p(y_* | \mathbf{x}_*, \mathbf{w}, \beta) p(\mathbf{w} | \mathbf{X}, \mathbf{y}, \beta, \lambda) d\mathbf{w}$$

Only  $\mathbf{w}$  is unknown with a posterior distribution so only  $\mathbf{w}$  has to be integrated out

$\mathcal{N}(y_* | \mathbf{w}^\top \mathbf{x}_*, \beta^{-1})$

$\mathcal{N}(\mathbf{w} | \boldsymbol{\mu}_N, \boldsymbol{\Sigma}_N)$

- The above is the marginalization of  $\mathbf{w}$  from  $\mathcal{N}(y_* | \mathbf{w}^\top \mathbf{x}_*, \beta^{-1})$ . Using LGM results

$$p(y_* | \mathbf{x}_*, \mathbf{X}, \mathbf{y}, \beta, \lambda) = \mathcal{N}(\boldsymbol{\mu}_N^\top \mathbf{x}_*, \beta^{-1} + \mathbf{x}_*^\top \boldsymbol{\Sigma}_N \mathbf{x}_*)$$

Can also derive it by writing  $y_* = \mathbf{w}^\top \mathbf{x}_* + \epsilon$  where  $\mathbf{w} \sim \mathcal{N}(\boldsymbol{\mu}_N, \boldsymbol{\Sigma}_N)$  and  $\epsilon \sim \mathcal{N}(0, \beta^{-1})$

- So we have a predictive mean  $\boldsymbol{\mu}_N^\top \mathbf{x}_*$  as well as an input-specific predictive variance
- In contrast, MLE and MAP make “plug-in” predictions (using the point estimate of  $\mathbf{w}$ )


$$\begin{aligned} p(y_* | \mathbf{x}_*, \mathbf{w}_{MLE}) &= \mathcal{N}(\mathbf{w}_{MLE}^\top \mathbf{x}_*, \beta^{-1}) && \text{- MLE prediction} \\ p(y_* | \mathbf{x}_*, \mathbf{w}_{MAP}) &= \mathcal{N}(\mathbf{w}_{MAP}^\top \mathbf{x}_*, \beta^{-1}) && \text{- MAP prediction} \end{aligned}$$

Since PPD also takes into account the uncertainty in  $\mathbf{w}$ , the predictive variance is larger

- Unlike MLE/MAP, variance of  $y_*$  also depends on the input  $\mathbf{x}_*$  (this, as we will see later, will be very useful in sequential decision-making problems such as active learning)

# Posterior Predictive Distribution: An Illustration


- Black dots are training examples



- Width of the shaded region at any  $x$  denotes the predictive uncertainty at that  $x$  ( $\pm$  one std-dev)
- Regions with more training examples have smaller predictive variance



# Nonlinear Regression



- Can extend the linear regression model to handle nonlinear regression problems
- One way is to replace the feature vectors  $\mathbf{x}$  by a nonlinear mapping  $\phi(\mathbf{x})$

$$p(y|\mathbf{x}, \mathbf{w}) = \mathcal{N}(\mathbf{w}^\top \phi(\mathbf{x}), \beta^{-1})$$

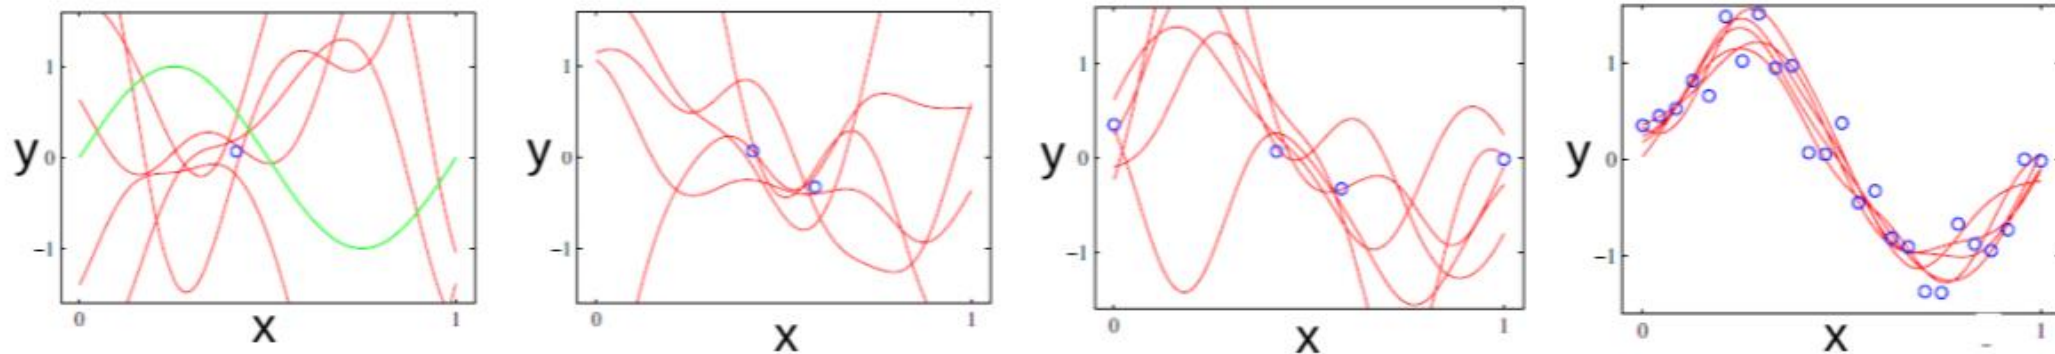
Can be pre-defined (e.g., replace a scalar  $x$  by polynomial mapping  $[1, x, x^2]$ ) or extracted by a pretrained deep neural net

- Alternatively, a [kernel function](#) can be used to implicitly define the nonlinear mapping
- More on nonlinear regression when we discuss [Gaussian Processes](#)




# More on Visualization of Uncertainty

- Figures below: Green curve is the true function and blue circles are observations
- Posterior of the nonlinear regression model: Some curves drawn from the posterior



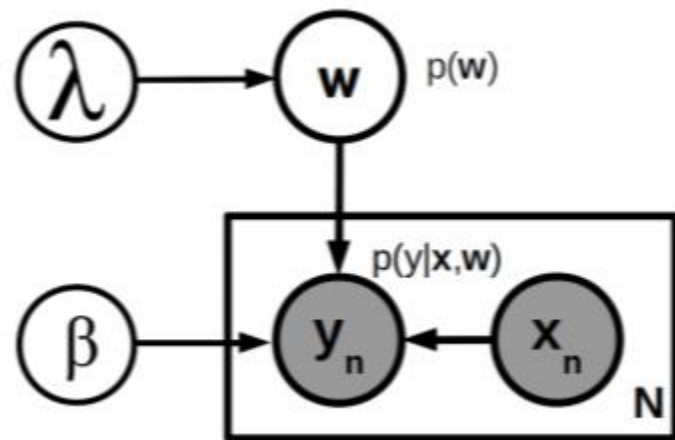
- PPD: Red curve is predictive mean, shaded region denotes predictive uncertainty





# Estimating Hyperparameters via MLE-II

- The probabilistic linear reg. model we saw had two hyperparams  $(\beta, \lambda)$ 
  - Thus total three unknowns  $(\mathbf{w}, \beta, \lambda)$



Need posterior over all the 3 unknowns

$$p(\mathbf{w}, \beta, \lambda | \mathbf{X}, \mathbf{y}) = \frac{p(\mathbf{y} | \mathbf{X}, \mathbf{w}, \beta, \lambda) p(\mathbf{w}, \lambda, \beta)}{p(\mathbf{y} | \mathbf{X})}$$

PPD would require integrating out all 3 unknowns

$$p(\mathbf{w}, \beta, \lambda | \mathbf{X}, \mathbf{y}) = \frac{p(\mathbf{y} | \mathbf{X}, \mathbf{w}, \beta, \lambda) p(\mathbf{w} | \lambda) p(\beta) p(\lambda)}{\int p(\mathbf{y} | \mathbf{X}, \mathbf{w}, \beta) p(\mathbf{w} | \lambda) p(\beta) p(\lambda) d\mathbf{w} d\lambda d\beta}$$

$$p(y_* | \mathbf{x}_*, \mathbf{X}, \mathbf{y}) = \int p(y_* | \mathbf{x}_*, \mathbf{w}, \beta) p(\mathbf{w}, \beta, \lambda | \mathbf{X}, \mathbf{y}) d\mathbf{w} d\beta d\lambda$$

- Posterior and PPD computation is intractable.
- If we just want point estimates for  $(\beta, \lambda)$  then MLE-II is an option

Called "MLE-II" because we are maximizing **marginal likelihood**, not the likelihood

And then compute  $p(\mathbf{w} | \mathbf{X}, \mathbf{y}, \hat{\beta}, \hat{\lambda})$  treating  $\hat{\beta}, \hat{\lambda}$  as given

$$(\hat{\beta}, \hat{\lambda}) = \operatorname{argmax}_{\beta, \lambda} \log p(\mathbf{y} | \mathbf{X}, \beta, \lambda)$$

For regression with Gaussian likelihood and Gaussian prior on  $\mathbf{w}$ , the marginal likelihood has an exact expression

Will see various other methods like EM, variational inference, MCMC, etc later

# Prob. Linear Regression: Some Other Variations

- Can use other likelihoods  $p(y_n | \mathbf{x}_n, \mathbf{w})$  and/or prior distribution  $p(\mathbf{w})$

- Laplace distribution for the likelihood

$$p(y_n | \mathbf{x}_n, \mathbf{w}) = \text{Lap}(y_n | \mathbf{w}^\top \mathbf{x}_n, b)$$

- Heteroskedastic noise in the likelihood, e.g.,

$$p(y_n | \mathbf{x}_n, \mathbf{w}) = \mathcal{N}(y_n | \mathbf{w}^\top \mathbf{x}_n, \beta_n^{-1})$$

Can even assume  $\beta_n$  to depend on input  $\mathbf{x}_n$


Different noise distribution  $\mathcal{N}(0, \beta_n^{-1})$  for each  $y_n$

- Feature-specific variances in the prior for  $\mathbf{w}$

$$p(\mathbf{w}) = \prod_{d=1}^D \mathcal{N}(w_d | 0, \lambda_d^{-1}) = \mathcal{N}(\mathbf{w} | \mathbf{0}, \Lambda^{-1})$$

This has the effect of having feature-specific regularization

Since we can also learn these precisions (e.g., using MLE-II), using such a prior, we can learn the importance of different features (**feature selection**) which isn't possible with a  $\mathcal{N}(\mathbf{w} | \mathbf{0}, \lambda^{-1} \mathbf{I})$  prior with spherical covariance



Diagonal precision/covariance matrix with  $\lambda_d$ 's along the columns of  $\Lambda$

